

Writing A Security Protocol On Top of NFC

Andrey Tydnyuk
Carnegie Mellon University
ait@andrew.cmu.edu

Srinivasan Seshan
Carnegie Mellon University
srini@cmu.edu

ABSTRACT

Near Field Communication technology, or NFC as it is commonly called, is a wireless radio communication standard that can be used to exchange information between a device and a tag at close range. It is a lot like RFID, but where RFID tags can be scanned at a distance of meters, NFC tags need to be scanned at a distance of centimeters. The upshot of this is that NFC allows two-way communication as opposed to RFID tags that generally only support reading from the tag¹.

In this paper we explore the possibility of adding a security protocol to NFC tags to allow secure location confirmation. For example, if a restaurant wants to make sure that only people that have been to it can submit reviews, it can place an NFC tag at the entrance/exit and confirm that customers have scanned the tag before they are allowed to submit a review online. We explore several different tag-device setups under several attack scenarios. We show that it is possible to develop a secure protocol with an on-phone TPM and a dumb tag that does not do computation. If we relax the TPM requirement, we can have a “good enough” solution that relies on the fact that only a certain fraction of users are malicious.

Categories and Subject Descriptors

H.4 [Mobile Systems Applications]: NFC, Security, Wireless

General Terms

Security, Reliability

Keywords

NFC, Wireless, Security, TPM, Location Services, Simulation

¹Newer RFID tags actually allow writing, but this is new and experimental

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

1. INTRODUCTION

NFC tag usage has been increasing since their mass production release. Tags are pretty cheap and versatile. Some popular uses of tags are to do location services such as checking in to places, or to change phone settings based on predefined rules. Some examples of the former are tags that are put up in front of restaurants or other commercial places that give you a code to a website to confirm that you have visited their location. This can be tied into different perks for the user such as discounts or sales. NFC tags can also be used for games or competitions that involve real-world locations. This can range from scavenger hunts that depend on the user scanning tags to get clues to games that assign territory based on number of visits to a particular location.

All of these uses rely on the fact that there won't be truly malicious users. Tolerating a few fake requests submitted through the web, without ever visiting the tag is fine, but being protected against DOS and other attacks is another. If the NFC tag associated with the service supports writing, then a malicious user can write a different clue or password or URL to the tag and thereby deny service or worse, redirect service to his own malicious page. The only way to fix this is for someone in charge to visit the site and write the correct password back to the tag. This requires a human intervention for every malicious user, and is not a reasonable solution.

On the other hand, if we disable writing and use a read-only tag we have a scenario where the password never changes. If the password never changes it can be put up online and accessed by anyone. While this might not be a worry for a local scavenger hunt that operates under an honor system, a review service operating in this scenario is rendered completely useless. If the tag was put in place to make sure that the reviews written are by actual past customers, the fact that you can put the password up online and get multiple uses out of it voids that guarantee.

In light of these difficulties we present several different systems for dealing with attacks against an NFC tag based service. The protocols and assumptions may differ for each setup, but there are several commonalities. Each setup consists of three parts. There is a tag that is located at a specific location that we want to securely transmit, there is a device (phone) that will be scanning the tag, and there is a web-server that will keep track of all of the scans and be in charge of accepting or rejecting requests.

We start our analysis by detailing several different attack models. Since there are a lot of attacks that can be done against this kind of service we cannot cover all of them, but

we list what we believe to be the easiest to execute and the most debilitating for the service itself. We will briefly mention other attacks that we did not consider in our analysis and give reasoning for our choice.

After going over the different attacks we consider two different setups to protect our service. The first setup assumes an on-phone TPM that can be used to shield the user from accessing certain information. This setup is found to be very effective at preventing malicious users from affecting the service. With our protocol we manage to prevent all but one of the attacks that we detailed in the previous section and we do not drop any requests from honest users. This is the best possible result, but it requires an on-phone TPM which isn't something that is commercially available.

For the second setup, we relax the on-phone TPM requirement and consider the lowest cost case of an average phone with NFC reading/writing capabilities and a dumb tag that cannot do any computation. In this case, we find that the protocol is vulnerable to several attacks. Our solution to alleviate this problem is the addition of a reputation system. We propose that the website keep a reputation record for each user and show that this approach helps us keep malicious users out of the system with minimal impact to honest users as long as the number of malicious users is about a tenth of the number of honest users.

At the end of the paper we go over some opportunities for further work.

2. ATTACK MODELS

In this section we will consider the different attacks that our service will have to deal with. We will consider attacks in order based on the ease of execution.

2.1 Bogus Submissions

Since the webserver is in charge of separating the valid requests from the invalid ones, the easiest attack to execute is to send bogus requests to the webserver. This does not require ever being next to the tag and can be done by anyone from the comfort of their own house.

The attack model that we consider for this is a maximum of 1 request per second. While it might seem strange that we assume such a low rate, especially considering that scripts can pump out hundreds of requests per second from a single machine, dealing with DDOS is out of the scope of this paper. Since there is no way a legitimate user can scan more than one tag per second, it seems reasonable to limit the number of requests that can come from a single IP to one per second.

2.2 Publicizing the Password

This attack is arguably simpler than the last one mentioned, but it requires going to the actual location of the tag and scanning it. Once the tag is scanned, the content of the tag is put up online for anyone to access.

The attack model for this is fairly simple as well. Since you actually need one person to go to the location, this attack is only harmful if the password has multiple uses. If it doesn't, this attack falls under another category that we will detail later.

2.3 Multiple Scan

This attack also involves a malicious user going to the location, but this time instead of just scanning the tag once,

the user scans the tag multiple times. He then potentially walks away with several passwords that he can distribute online at will.

For this attack, we accept that the passwords may be single-use, but the attacker can scan the tag an unlimited number of times if necessary and put the list up online. This way we get a similar effect to the **Publicized Password** attack, with the additional limitation that it takes more time for the user to stand and scan the tag several times. Since malicious users don't really need a reason to be malicious, we don't really regard it as a real limitation and just assume that there will be someone who is willing to stand at the tag for hours recording thousands of passwords.

2.4 Tag Modification

This is also an attack that involves a malicious user going to the location of the tag. In this scenario, the user writes garbage or malicious data to the tag.

For this attack we assume that the attacker can write anything to the tag and the frequency with which this happens is proportionate to the number of malicious users in the system. We assume that this kind of attack does not happen that often due to the requirement that the malicious user has to be in range of the tag itself.

2.5 Man in the Middle

This attack requires the most effort on the attacker's part, and is also the most difficult to defend against. This attack involves an attacker scanning the tag, and instead of submitting the password, forwarding it to a friend. The rest of the protocol proceeds a fashion such that all communication between phone and webserver goes through the middleman that is pretending to be at the tag location.

2.6 Other Attacks

We never cover snooping attacks for the simple reason that NFC communication operates at a close range, so it is difficult to snoop on a communication between a phone and a tag without being painfully obvious to the honest user. We also never consider attacks on other layers of the network stack because that is out of the scope of this paper.

3. PROTOCOL 1: PHONE WITH TPM + DUMB TAG

3.1 Outline

This is a quick rundown of the protocol:

1. The device reads the tag and the tag sends the device a code, k
2. The key, k , goes into sealed storage
3. The device sends the phone IMEI and k to the webserver
4. The webserver checks k against the key that it is expecting
5. The key matches so the webserver generates the next key, l , to write to the tag
6. The webserver records the fact that the device's request was accepted

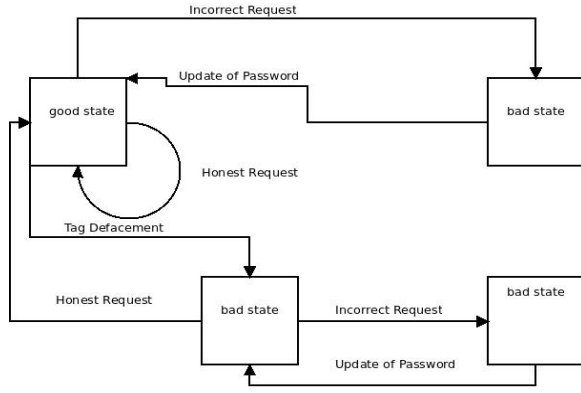


Figure 1: State diagram for NFC tag protocol

7. The ACCEPTED message and l are sent to the device
8. The device writes l to the tag so the next person can use it

That was a step by step overview of what happens when an honest user uses the protocol. The interesting stuff happens when malicious users enter the picture. For a comprehensive look at this we need to look at the different states that the protocol can be in depending on what requests it gets.

3.2 Protocol States

Figure 1 contains the state diagram that we have to consider. This diagram details the different ways that users can interact with the system and also shows the requests that should be accepted and rejected. There are three possible transitions between states. One is an honest user interacting with the system, one is a malicious user sending requests to the webserver directly, and the third is a malicious user defacing the NFC tag.

From the description we can see that any request that moves the state of the system from good to bad should be rejected because it is malicious. Likewise, any request that moves the state from bad to good should be accepted because it is an honest user.

3.3 Argument for Correctness

We will refer to Figure 1 and argue that the protocol that we have described above guarantees all that we have described.

If the tag is in a good state i.e. it contains the next password that the webserver is expecting, we can see that a sequence of honest users will not result to a change to a bad state. Each user will read the tag, submit the correct code to the server and then write the next correct code to the tag. However, if one of these honest users drags their phone away too fast, or if a malicious user comes in and defaces the tag, we transition to a bad state.

From this state, if an honest user comes in afterwards he will read the wrong key and his request will get rejected. To deal with this, we rely on the next honest request to fix our mistake. When a valid request is accepted, the webserver goes back and checks to make sure that the person that wrote that message to the tag gets their request accepted as well. This way, even though it looks like an honest user has to sacrifice a transaction to fix the state, we can actually

Table 1: Attack Prevention Table for Protocol 1

Attack	Protocol 1	Protocol 2
Bogus Submissions	YES	YES
Publicizing the Password	YES	YES
Multiple Scan	YES	YES
Tag Modification	YES	Vote Lost
Man In The Middle	YES	NO

roll back their rejection and approve it provided that another valid request comes in after them. It is worth noting that this only works with the TPM because otherwise malicious users can just pass codes around and confirm each other's requests.

Now that we have covered two of the states we can look at the cases where requests get submitted from the web. If a request is submitted from the web, we assume that it is not going to be the right password since the chances of that happening are 1 in 2^{15} in our proof of concept implementation, and can be arbitrarily lower if we choose a longer cryptographically random key. If an incorrect request is sent from any state, the transition is to a bad state and the webserver will reject the request. After each transition to a bad state, the state will then switch back to the one it was at previously after a new password is generated. This guarantees that no web requests will affect the protocol.

3.4 Attack Prevention

Please refer to Figure 2 to see how our protocol fares against attacks. We can deal with bogus submissions as we argued with the state model. Each password has one use, so publicizing it has no real benefit. Multiple scans are disallowed by the web interface since the phone has to submit an IMEI with each request, and even if that is somehow avoided the requests are tied to the user scanning the tag so they cannot be passed to a third party. Tag modification is fixable with one honest user, and then that honest user can get their request accepted when another honest user scans the tag after them. A MITM attack can be prevented with the use of a TPM to hide the password from the user, thus preventing him from forwarding it to anyone.

3.5 Implementation

We have implemented a proof of concept for this protocol, but since we did not have a phone with a TPM on it we just worked under the assumption that the passwords that the web interface gave the users would be inaccessible. We tested the protocol with malicious and honest requests and showed that the state transitions that we expected were reproduceable in the real world.

4. PROTOCOL 2: PHONE WITHOUT TPM + DUMB TAG

4.1 Outline

This is a quick rundown of the protocol:

1. The device reads the tag and the tag sends the device a code, k
2. The device sends the phone IMEI and k to the webserver

3. The webserver checks k against the key that it is expecting
4. The key matches so the webserver generates the next key, l , to write to the tag
5. The webserver records the fact that the device's request was accepted
6. The ACCEPTED message and l are sent to the device
7. The device writes l to the tag so the next person can use it

If you compare this protocol to the one that we just considered, you can see that there are almost no differences. The protocol is practically the same, but now we have relaxed the requirement of an on-phone TPM. This is very significant, because it means that malicious users have unfettered access to the keys sent to them by the webserver and tag.

4.2 Attack Prevention

For this protocol we cannot really show correctness because we don't actually have it. What we do have is a very good system for minimizing the possible damage that the attacks can do.

Please refer to Figure 2 to see how our protocol fares against attacks. We can deal with bogus submissions the same way that we did in the first protocol. In this regard, nothing has changed. False requests have no real impact. Each password has one use, so publicizing it has no real benefit. Multiple scans are disallowed by the web interface since the phone has to submit an IMEI with each request, and even if that is somehow avoided the requests are tied to the user scanning the tag so they cannot be passed to a third party en masse. What can happen is that it can be passed a single time (a man in the middle attack that we will discuss below). Tag modification is fixable with one honest user, but in this case we cannot roll the log back and re-accept the user that fixed the tag. This is because we have no guarantee that two malicious users won't get together and confirm each other's requests. All it takes in this scenario is for one to submit a faulty request, and then pass the new password that he is given to his friend. That way both of them now have valid entries in the log. Because of this, we have to turn this feature off and deal with the fact that every time a malicious user modifies the tag, we will use one honest user's submission. We can implement a clever mechanism to deal with this to prevent it from becoming a real problem. Unfortunately, a MITM attack has full reign under this protocol because we have no means to prevent people from passing their keys to other users. We argue that this isn't really all that bad considering the fact that at least someone has to be at the location and that MITM requires some coordination because it has to be completed before another user walks up to the tag and scans it. Since there isn't a good way to prevent MITM with a dumb tag and a phone that can't shield information from the user, we will focus on cutting down the number of malicious users in the system in general to lessen their effect on the good users.

4.3 Reputation System

A reputation system seems like the obvious choice for this scenario. If users have to authenticate to the webserver and

Table 2: Simulations of Protocol 2

Total Requests	Bad Users	Reputation Enabled	% Dropped
3058	100	NO	10.16
2976	100	YES	9.22
2187	10	YES	1.72

if their acceptances and rejections are tied to their account, it makes it harder for repeat offenders to affect the system. Privacy is a concern, but in this case the only thing that the reputation system has to tie to the user is the number of times their requests have been accepted or rejected. We do not need to know where exactly these requests came from, so it can't really be tied to a real identity.

The way that we propose to do the reputation system involves adding certain numbers of strikes to users for actions that are possibly malicious. These actions are associated with being one of the people around the transitions to the bad states in the state diagram in Figure 1. If the same user is seen transitioning the system to a bad state it is safe to assume that he is a malicious user, and then he can be banned from the system. This prevents outlying malicious users from flooding the system with bad requests and repeatedly defacing the tag, which as we covered above, results in the loss of good votes.

Since it is hard to tell how a reputation system will work without trying it first, we ran several simulations.

4.4 Simulations

We ran the simulations in a similar way. We had several threads for malicious users, and a thread for honest users. We ran each simulation for 30 simulated days and collected data at the end for the number of good requests that were dropped from the log. The variables that we tweaked in the simulation were the ratio of malicious users to good users and the number of requests issued to the service. These variables were reflected in the timing of the threads. For example, in a system that is very popular and with a small fraction of malicious users we have the good thread printing often, and the malicious threads printing rarely. We then parse the resulting logfile and compare the number of confirmed accepted requests to the expected number. We return the result as a percentage.

When we ran the simulations without a reputation system, the number of honest user requests that was dropped was pretty much equal to the number of tag defacements that occurred during the simulation. This makes intuitive sense because every time a tag is brought out of sync, a good user has to sacrifice a vote in order to sync it up again. With the addition of a reputation system, however, the numbers that we got were always lower than without the system. Especially when the malicious users were small in number, they got weeded out of the system fairly quickly. This bodes well for the system because since the users have to be at the physical location, it prevents several users that live or work in the vicinity from strongly affecting honest users.

We did not take into account malicious requests sent from the web because they would not affect our protocol either way. With the reputation system however, they would actually improve results and allow us to stop malicious users even faster. Please refer to Table 2 for the results of the simulations:

5. RESULTS

We can see from the results of the simulation that the second protocol works fairly well even without a TPM. This is the lowest cost solution that we can hope for, and with the addition of a reputation system we can have a fairly stable system that is relatively resistant to attacks by malicious users.

Of course the better alternative is to use the first protocol where we don't have to bother with a reputation system and don't have to worry about MITM attacks, but no commercial phone currently has a TPM on it. It might be a reasonable thing to expect in the future since a lot of laptops are currently being released with TPMs built in, but it is not a reasonable thing to expect in the next year.

6. FURTHER RESEARCH

This paper focused heavily on the webserver and the device side of the problem. If we consider tags that have timing code or can cryptographically sign messages we can consider other possible solutions. Dumb tags are the most common form right now and will definitely be cheaper than the smarter alternative, so any protocol that assumes a smart tag has to look to the future in terms of real-world implementation.